

Redes Neuronales

Ejercicio 1:

```
1  import numpy as np # type: ignore
2  import pandas as pd # type: ignore
3  from sklearn.neural_network import MLPRegressor # type: ignore
4  from sklearn.model_selection import KFold # type: ignore
5  from sklearn.metrics import mean_squared_error, r2_score # type: ignore
6
7  # Cargar el archivo CSV
8  data = pd.read_csv('crime_data.csv')
9
10
11  x = data[['M', 'W', 'S', 'P']].values
12  y = data['MR'].values
13  reg = MLPRegressor(hidden_layer_sizes=(10, 10), max_iter=10000)
14  reg.fit(x, y)
15
16  kf = KFold(n_splits=5, shuffle=True)
17
18  cv_y_test = []
19  cv_y_pred = []
20
21  for train_index, test_index in kf.split(x, y):
22      # Fase de entrenamiento
23      x_train = x[train_index, :]
24      y_train = y[train_index]
25
26      reg_i = MLPRegressor(hidden_layer_sizes=(10, 10), max_iter=10000)
27      reg_i.fit(x_train, y_train)
28
29      # Fase de prueba
30      x_test = x[test_index, :]
31      y_test = y[test_index]
32      y_pred = reg_i.predict(x_test)
33
34      cv_y_test.append(y_test)
35      cv_y_pred.append(y_pred)
36
37  # Calcular el error cuadrático medio y el coeficiente de determinación R2
38  y_test_concat = np.concatenate(cv_y_test)
39  y_pred_concat = np.concatenate(cv_y_pred)
40
41  print("MSE:", mean_squared_error(y_test_concat, y_pred_concat))
42  print("R2:", r2_score(y_test_concat, y_pred_concat))
```

```
MSE: 67.99105879995845
R2: 0.39624763172949773
```

¿Consideras que el modelo perceptrón multicapa es efectivo para modelar los datos del problema? ¿Por qué? No, tiene un error cuadrado medio de 67, el cual es excesivo, y el R2 nos da 0.396, bastante lejos de 1 lo cual nos da un nivel bajo de exactitud.

¿Qué modelo es mejor para los datos de criminalidad, el lineal o el perceptrón multicapa? ¿Por qué? Los datos parecen ser lineales, por lo que un perceptrón multicapa no es el indicado para este tipo de problemas.

Ejercicio 2:

```
1  import numpy as np # type: ignore
2  import pandas as pd # type: ignore
3  from sklearn.neural_network import MLPRegressor # type: ignore
4  from sklearn.model_selection import GridSearchCV, KFold # type: ignore
5  from sklearn.metrics import mean_squared_error, r2_score # type: ignore
6
7  data = pd.read_csv('M_4.txt', delim_whitespace=True, header=None)
8
9  x = data.iloc[:, 2:].values
10 y = data.iloc[:, 0].values
11
12 param_grid = {
13     'hidden_layer_sizes': [(20,) * 5],
14     'activation': ['relu', 'tanh', 'logistic'],
15     'solver': ['adam', 'sgd'],
16     'alpha': [0.0001, 0.0002, 0.0003],
17     'max_iter': [10000, 20000, 30000]
18 }
19
20 mlp = MLPRegressor()
21
22 kf = KFold(n_splits=5, shuffle=True, random_state=42)
23 grid_search = GridSearchCV(estimator=mlp, param_grid=param_grid, cv=kf, scoring='neg_mean_squared_error', n_jobs=-1)
24
25 grid_search.fit(x, y)
26
27 print("Mejores parámetros encontrados:")
28 best_params = grid_search.best_params_
29 print(best_params)
30
31 best_model = grid_search.best_estimator_
32
33 cv_results = grid_search.cv_results_
34 mean_test_scores = cv_results['mean_test_score']
35 std_test_scores = cv_results['std_test_score']
36
37 print("\nResultados de la búsqueda:")
38 for mean_score, std_score, params in zip(mean_test_scores, std_test_scores, cv_results['params']):
39     print(f"Parámetros: {params} | MSE: {-mean_score:.4f} (+/- {std_score:.4f})")
40
41 y_pred = best_model.predict(x)
42 print("\nEvaluación del mejor modelo:")
43 print("MSE:", mean_squared_error(y, y_pred))
44 print("R2:", r2_score(y, y_pred))
45
46 print(f"Parámetros óptimos: {best_params}")
```

```

Resultados de la búsqueda:
Parámetros: {'activation': 'relu', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'adam'} | MSE: 1.0010 (+/- 0.0071)
Parámetros: {'activation': 'relu', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'sgd'} | MSE: 1.1508 (+/- 0.1614)
Parámetros: {'activation': 'relu', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'adam'} | MSE: 0.9990 (+/- 0.0069)
Parámetros: {'activation': 'relu', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'sgd'} | MSE: 1.0712 (+/- 0.0915)
Parámetros: {'activation': 'relu', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'adam'} | MSE: 1.0606 (+/- 0.1667)
Parámetros: {'activation': 'relu', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'sgd'} | MSE: 1.1074 (+/- 0.2266)
Parámetros: {'activation': 'relu', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'adam'} | MSE: 1.0590 (+/- 0.1841)
Parámetros: {'activation': 'relu', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'sgd'} | MSE: 1.1447 (+/- 0.0874)
Parámetros: {'activation': 'relu', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'adam'} | MSE: 1.1183 (+/- 0.1908)
Parámetros: {'activation': 'relu', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'sgd'} | MSE: 0.9804 (+/- 0.0198)
Parámetros: {'activation': 'relu', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'adam'} | MSE: 1.0574 (+/- 0.1965)
Parámetros: {'activation': 'relu', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'sgd'} | MSE: 1.0918 (+/- 0.0815)
Parámetros: {'activation': 'relu', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'adam'} | MSE: 1.1144 (+/- 0.2271)
Parámetros: {'activation': 'relu', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'sgd'} | MSE: 1.0944 (+/- 0.0937)
Parámetros: {'activation': 'relu', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'adam'} | MSE: 1.0166 (+/- 0.0737)
Parámetros: {'activation': 'relu', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'sgd'} | MSE: 1.1819 (+/- 0.1660)
Parámetros: {'activation': 'relu', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'adam'} | MSE: 0.9657 (+/- 0.1487)
Parámetros: {'activation': 'relu', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'sgd'} | MSE: 1.1477 (+/- 0.0808)
Parámetros: {'activation': 'tanh', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'adam'} | MSE: 0.9694 (+/- 0.1818)
Parámetros: {'activation': 'tanh', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'sgd'} | MSE: 1.0232 (+/- 0.2747)
Parámetros: {'activation': 'tanh', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'adam'} | MSE: 1.1851 (+/- 0.1569)
Parámetros: {'activation': 'tanh', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'sgd'} | MSE: 1.2923 (+/- 0.1498)
Parámetros: {'activation': 'tanh', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'adam'} | MSE: 1.1276 (+/- 0.1982)
Parámetros: {'activation': 'tanh', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'sgd'} | MSE: 1.0865 (+/- 0.2370)
Parámetros: {'activation': 'tanh', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'adam'} | MSE: 1.0989 (+/- 0.1207)
Parámetros: {'activation': 'tanh', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'sgd'} | MSE: 1.1149 (+/- 0.2855)
Parámetros: {'activation': 'tanh', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'adam'} | MSE: 1.1487 (+/- 0.2563)
Parámetros: {'activation': 'tanh', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'sgd'} | MSE: 1.1143 (+/- 0.2257)
Parámetros: {'activation': 'tanh', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'adam'} | MSE: 1.2365 (+/- 0.3275)
Parámetros: {'activation': 'tanh', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'sgd'} | MSE: 1.0647 (+/- 0.1823)
Parámetros: {'activation': 'tanh', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'adam'} | MSE: 0.9550 (+/- 0.1781)
Parámetros: {'activation': 'tanh', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'sgd'} | MSE: 1.1227 (+/- 0.1078)
Parámetros: {'activation': 'tanh', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'adam'} | MSE: 0.9943 (+/- 0.1527)
Parámetros: {'activation': 'tanh', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'sgd'} | MSE: 1.1822 (+/- 0.2453)
Parámetros: {'activation': 'tanh', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'adam'} | MSE: 0.9776 (+/- 0.1127)
Parámetros: {'activation': 'tanh', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'sgd'} | MSE: 1.1725 (+/- 0.2782)
Parámetros: {'activation': 'logistic', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'adam'} | MSE: 1.6896 (+/- 1.1228)
Parámetros: {'activation': 'logistic', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'sgd'} | MSE: 4.0035 (+/- 0.2005)
Parámetros: {'activation': 'logistic', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'adam'} | MSE: 1.7273 (+/- 1.2166)
Parámetros: {'activation': 'logistic', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'sgd'} | MSE: 4.0073 (+/- 0.2033)
Parámetros: {'activation': 'logistic', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'adam'} | MSE: 2.2388 (+/- 1.4941)
Parámetros: {'activation': 'logistic', 'alpha': 0.0001, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'sgd'} | MSE: 4.0093 (+/- 0.2000)
Parámetros: {'activation': 'logistic', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'adam'} | MSE: 1.7046 (+/- 1.2056)
Parámetros: {'activation': 'logistic', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'sgd'} | MSE: 4.0083 (+/- 0.2033)
Parámetros: {'activation': 'logistic', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'adam'} | MSE: 1.0952 (+/- 0.1830)
Parámetros: {'activation': 'logistic', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'sgd'} | MSE: 4.0115 (+/- 0.2061)
Parámetros: {'activation': 'logistic', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'adam'} | MSE: 2.7880 (+/- 1.5592)
Parámetros: {'activation': 'logistic', 'alpha': 0.0002, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'sgd'} | MSE: 4.0027 (+/- 0.2006)
Parámetros: {'activation': 'logistic', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'adam'} | MSE: 2.1860 (+/- 1.5217)
Parámetros: {'activation': 'logistic', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'sgd'} | MSE: 4.0027 (+/- 0.1983)
Parámetros: {'activation': 'logistic', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 20000, 'solver': 'adam'} | MSE: 1.2119 (+/- 0.2458)
Parámetros: {'activation': 'logistic', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 30000, 'solver': 'sgd'} | MSE: 4.0098 (+/- 0.2051)

Evaluación del mejor modelo:
MSE: 0.34093139706635567
R2: 0.9147671507334111
Parámetros óptimos: {'activation': 'tanh', 'alpha': 0.0003, 'hidden_layer_sizes': (20, 20, 20, 20, 20), 'max_iter': 10000, 'solver': 'adam'}

```

¿Observas alguna mejora importante al optimizar el tamaño de la red? ¿Es el resultado que esperabas? Argumenta tu respuesta.

Dependiendo de los parámetros podemos observar MSEs que varían de 0.34 hasta 4, por lo que el encontrar los valores y parámetros óptimos nos da mucho mejores resultados.

¿Qué inconvenientes hay al encontrar el tamaño óptimo de la red? ¿Por qué?

Al crecer la red neuronal estamos añadiendo tiempo computacional, el cual en una computadora de escritorio, incluso de gama alta, puede convertirse en un problema. No solo esto si no que al crecer más la red neuronal podría caer en un sobreajuste a los datos.

Ejercicio 3:

```
1  import numpy as np # type: ignore
2  import pandas as pd # type: ignore
3  import torch # type: ignore
4  import torch.nn as nn # type: ignore
5  import torch.optim as optim # type: ignore
6  from sklearn import datasets # type: ignore
7  from sklearn.model_selection import StratifiedKFold # type: ignore
8  from sklearn.metrics import accuracy_score # type: ignore
9  import matplotlib.pyplot as plt # type: ignore
10
11  data = pd.read_csv('P1_4.txt', delim_whitespace=True, header=None)
12
13  x = data.iloc[:, 2:].values
14  y = data.iloc[:, 0].values
15
16  unique_classes = np.unique(y)
17  class_map = {cls: idx for idx, cls in enumerate(unique_classes)}
18  y = np.array([class_map[cls] for cls in y])
19
20  x = torch.tensor(x, dtype=torch.float32)
21  y = torch.tensor(y, dtype=torch.long)
22
23  class SingleNeuronPerceptron(nn.Module):
24      def __init__(self, input_dim, output_dim):
25          super(SingleNeuronPerceptron, self).__init__()
26          self.fc = nn.Linear(input_dim, output_dim)
27
28      def forward(self, x):
29          return self.fc(x)
30
31  loss_function = nn.CrossEntropyLoss()
32
33  num_epochs = 1000
34  learning_rate = 0.01
35  input_dim = x.shape[1]
36  output_dim = len(np.unique(y.numpy()))
37
38  accuracy_list = []
39
```

```
kf = StratifiedKFold(n_splits=5, shuffle=True)

for fold, (train_index, test_index) in enumerate(kf.split(x, y), 1):
    x_train = x[train_index, :]
    y_train = y[train_index]
    x_test = x[test_index, :]
    y_test = y[test_index]

    model = SingleNeuronPerceptron(input_dim, output_dim)
    optimizer = optim.SGD(model.parameters(), lr=learning_rate)


    # Training phase
    for epoch in range(num_epochs):
        model.train()
        outputs = model(x_train)
        loss = loss_function(outputs, y_train)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Calculate accuracy
        model.eval()
        with torch.no_grad():
            test_outputs = model(x_test)
            _, y_pred = torch.max(test_outputs.data, 1)
            accuracy = accuracy_score(y_test.numpy(), y_pred.numpy())
            accuracy_list.append(accuracy)

        # Print results
        if (epoch+1) % 10 == 0:
            print(f"Fold {fold}, Epoch {epoch+1}, Loss: {loss.item()}, Accuracy: {accuracy}")

epochs = range(1, len(accuracy_list) + 1)
plt.plot(epochs, accuracy_list, 'b-', label='Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Epoch vs Accuracy')
plt.legend()
plt.show()
```



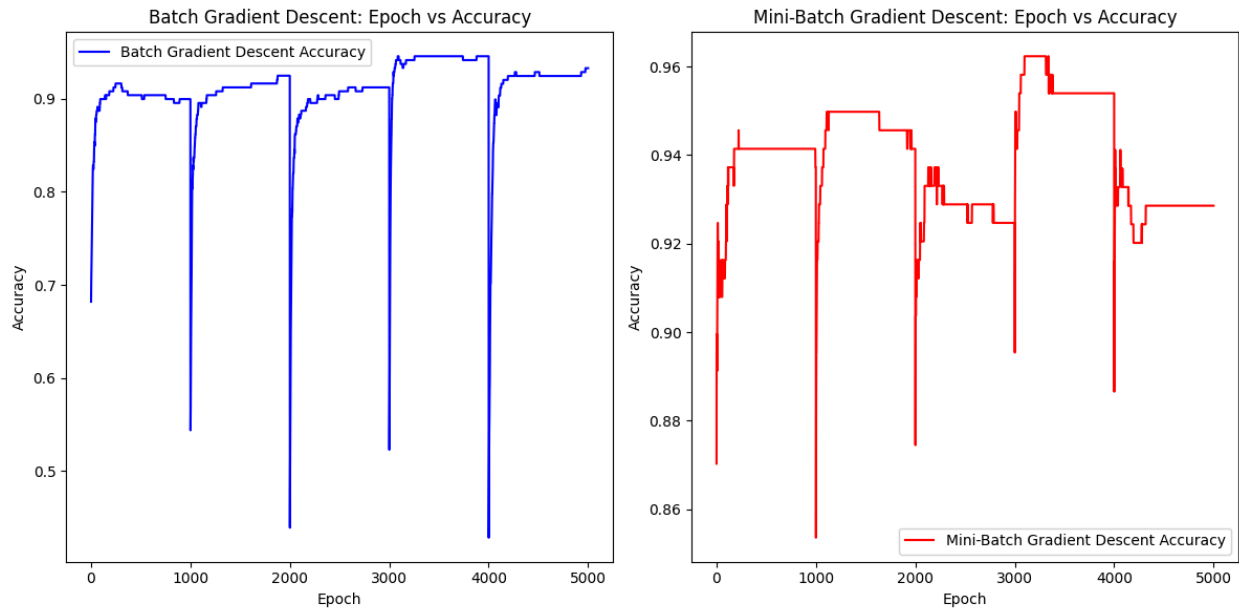
```

1  import numpy as np # type: ignore
2  import pandas as pd # type: ignore
3  import torch # type: ignore
4  import torch.nn as nn # type: ignore
5  import torch.optim as optim # type: ignore
6  from sklearn.model_selection import StratifiedKFold # type: ignore
7  from sklearn.metrics import accuracy_score # type: ignore
8  import matplotlib.pyplot as plt # type: ignore
9  from torch.utils.data import DataLoader, TensorDataset # type: ignore
10
11 # Load data
12 data = pd.read_csv('P1_4.txt', delim_whitespace=True, header=None)
13
14 x = data.iloc[:, 2:].values
15 y = data.iloc[:, 0].values
16
17 unique_classes = np.unique(y)
18 class_map = {cls: idx for idx, cls in enumerate(unique_classes)}
19 y = np.array([class_map[cls] for cls in y])
20
21 x = torch.tensor(x, dtype=torch.float32)
22 y = torch.tensor(y, dtype=torch.long)
23
24 class SingleNeuronPerceptron(nn.Module):
25     def __init__(self, input_dim, output_dim):
26         super(SingleNeuronPerceptron, self).__init__()
27         self.fc = nn.Linear(input_dim, output_dim)
28
29     def forward(self, x):
30         return self.fc(x)
31
32 loss_function = nn.CrossEntropyLoss()
33
34 num_epochs = 1000
35 learning_rate = 0.01
36 batch_size = 32
37 input_dim = x.shape[1]
38 output_dim = len(unique_classes)
39
40 accuracy_list_batch = []
41 accuracy_list_mini_batch = []
42

```

```
43 kf = StratifiedKFold(n_splits=5, shuffle=True)
44
45 for fold, (train_index, test_index) in enumerate(kf.split(x, y), 1):
46     # Separate training and test data
47     x_train = x[train_index, :]
48     y_train = y[train_index]
49     x_test = x[test_index, :]
50     y_test = y[test_index]
51
52     train_dataset = TensorDataset(x_train, y_train)
53     train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
54
55     # Batch Gradient Descent
56     model = SingleNeuronPerceptron(input_dim, output_dim)
57     optimizer = optim.SGD(model.parameters(), lr=learning_rate)
58
59     # Training phase
60     for epoch in range(num_epochs):
61         model.train()
62         outputs = model(x_train)
63         loss = loss_function(outputs, y_train)
64
65         optimizer.zero_grad()
66         loss.backward()
67         optimizer.step()
68
69         # Calculate accuracy
70         model.eval()
71         with torch.no_grad():
72             test_outputs = model(x_test)
73             _, y_pred = torch.max(test_outputs.data, 1)
74             accuracy = accuracy_score(y_test.numpy(), y_pred.numpy())
75             accuracy_list_batch.append(accuracy)
76
77     # Print results
78     if (epoch+1) % 10 == 0:
79         print(f"Fold {fold}, Batch Gradient Epoch {epoch+1}, Loss: {loss.item()}, Accuracy: {accuracy}")
80
81     # Mini-Batch Gradient Descent
82     model = SingleNeuronPerceptron(input_dim, output_dim)
83     optimizer = optim.SGD(model.parameters(), lr=learning_rate)
84
```

```
85     # Training phase
86     for epoch in range(num_epochs):
87         model.train()
88         for batch_x, batch_y in train_loader:
89             outputs = model(batch_x)
90             loss = loss_function(outputs, batch_y)
91
92             optimizer.zero_grad()
93             loss.backward()
94             optimizer.step()
95
96         # Calculate accuracy
97         model.eval()
98         with torch.no_grad():
99             test_outputs = model(x_test)
100             _, y_pred = torch.max(test_outputs.data, 1)
101             accuracy = accuracy_score(y_test.numpy(), y_pred.numpy())
102             accuracy_list_mini_batch.append(accuracy)
103
104         # Print results
105         if (epoch+1) % 10 == 0:
106             print(f"Fold {fold}, Mini-Batch Gradient Epoch {epoch+1}, Loss: {loss.item()}, Accuracy: {accuracy}")
107
108     epochs_batch = range(1, len(accuracy_list_batch) + 1)
109     epochs_mini_batch = range(1, len(accuracy_list_mini_batch) + 1)
110
111     plt.figure(figsize=(12, 6))
112
113     plt.subplot(1, 2, 1)
114     plt.plot(epochs_batch, accuracy_list_batch, 'b-', label='Batch Gradient Descent Accuracy')
115     plt.xlabel('Epoch')
116     plt.ylabel('Accuracy')
117     plt.title('Batch Gradient Descent: Epoch vs Accuracy')
118     plt.legend()
119
120     plt.subplot(1, 2, 2)
121     plt.plot(epochs_mini_batch, accuracy_list_mini_batch, 'r-', label='Mini-Batch Gradient Descent Accuracy')
122     plt.xlabel('Epoch')
123     plt.ylabel('Accuracy')
124     plt.title('Mini-Batch Gradient Descent: Epoch vs Accuracy')
125     plt.legend()
126
127     plt.tight_layout()
128     plt.show()
```

¿El modelo de una neurona es suficiente para modelar el conjunto de datos de este problema?

Sí, después de varios epochs casi todas las precisiones eran de arriba de 90, llegando incluso hasta 0.96, por lo que los modelos hicieron un muy buen trabajo de predecir los resultados de nuestro dataset.

Link del código:

https://drive.google.com/file/d/1FwpCMlyvPV2p_DaThzclk1I8vDKWgoKa/view?usp=sharing